

Running Apache Spark in a Docker Container

The following instructions can be used to install Spark in a SLES 12 SP1/Docker instance running on IBM LinuxONE Community Cloud. One example provides an interactive Spark shell to store and read values from an array. A second example starts Spark master and worker nodes to estimate the value of Pi.

Prerequisites

Sign up for a LinuxONE Community Cloud trial account

If you have not done so already, register at <http://www.ibm.com/linuxone/try> for a 120-day trial account. You will receive an email containing credentials to access the LinuxONE Community Cloud self-service portal. This is where you can start exploring all our available services.

Deploy a virtual server instance

If you have not deployed a virtual server already, please follow these instructions <http://developer.ibm.com/linuxone/wp-content/uploads/sites/57/virtual-servers-quick-start.pdf> to create one before proceeding. Make sure you select a flavor (resource definition) with 4 GB of memory.

This quick-start guide has been tested with the following Linux distributions:

- SUSE Linux Enterprise Server (SLES) 12

Step 1: Create the Docker container

1. Switch to root user for this exercise.
> `sudo -i`

```
linux1@dockerspark2:~> sudo -i
dockerspark2:~ #
```

2. Start the Docker service on your Linux system.
`systemctl start docker`

```
dockerspark2:~ # systemctl start docker
dockerspark2:~ #
```

3. Download the Spark Docker file from the repository.
`docker pull brunswickheads/spark-1.5.2-s390x`

```
dockersparks2:/ # docker pull brunswickheads/spark-1.5.2-s390x
Using default tag: latest
latest: Pulling from brunswickheads/spark-1.5.2-s390x
b5257eca2dd8: Pull complete
78745e6c93f9: Pull complete
b6a646867400: Pull complete
b89ba56b8f32: Pull complete
ca77adadcce6: Pull complete
a8e5f3181b12: Pull complete
Digest: sha256:c55f9185e571fbbccd6a2a2c126aed51700a441c67c8368da24b060612c0be7a
Status: Downloaded newer image for brunswickheads/spark-1.5.2-s390x:latest
```

4. Launch a Spark Docker container and open the ports required for Spark. Start a bash shell within the container.

```
# docker run -p 8080:8080 -p 8081:8081 -p 7077:7077 -p 6066:6066 -i -t
brunswickheads/spark-1.5.2-s390x /bin/bash
```

```
dockersparks2:~ # docker run -p 8080:8080 -p 8081:8081 -p 7077:7077 -p 6066:6066 -i -t brunswickheads/spark-1.5.2-s390x /bin/bash
bash-4.2#
```

5. Exit the bash shell.

```
# exit
```

```
bash-4.2# exit
exit
dockersparks2:/ #
```

6. The docker container is stopped when you exit the bash shell.

```
# docker ps
```

```
dockersparks2:~ # docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES

7. Switch to the Docker containers folder.

```
# cd /data/docker/containers/
```

```
dockersparks2:/ # cd /data/docker/containers/
dockersparks2:/data/docker/containers #
```

8. List the contents of the folder to view your Docker containers.

```
# ls
```

```
dockersparks2:/data/docker/containers # ls
422caf42253098ad225e2cf38ba91c335f04be6bb60f8733941b0ead4db35710
dockersparks2:/data/docker/containers #
```

Step 2: Start the Apache Spark shell

1. Restart your Docker container using the first three characters of your Docker container directory.

docker start 422

```
dockersparks2:/data/docker/containers # docker start 422
422
dockersparks2:/data/docker/containers # █
```

2. Check the status of your running Docker container.

docker ps

Note your Docker container id.

```
dockersparks2:/data/docker/containers # docker ps
CONTAINER ID        IMAGE               NAMES             COMMAND           CREATED           STATUS            PORTS
422caf422538       brunswickheads/spark-1.5.2-s390x  "/bin/bash"      45 minutes ago   Up 6 minutes     0.0.0.0:6066->6066/tcp, 0.0.0.0:7077->7077/tcp
dockersparks2:/data/docker/containers # █
```

3. Attach to your running Docker container to restart the bash shell.

- a) # docker attach 422
- b) Hit 'enter' or 'return' again

```
dockersparks2:/data/docker/containers # docker attach 422
bash-4.2# █
```

4. Start an interactive Spark shell.

/opt/ibm/spark/bin/spark-shell

```
bash-4.2# /opt/ibm/spark/bin/spark-shell
log4j:WARN No appenders could be found for logger (org.apache.hadoop.metrics2.lib.MutableMetricsFactor
y).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Using Spark's repl log4j profile: org/apache/spark/log4j-defaults-repl.properties
To adjust logging level use sc.setLogLevel("INFO")
Welcome to

      ____ _ 
     |    _ \| | | |__|_|_||_|
     |   || \ V V / __/ __/| |
     |___||_\_/\/____\____\_|_| version 1.5.2

Using Scala version 2.10.4 (IBM J9 VM, Java 1.8.0)
Type in expressions to have them evaluated.
Type :help for more information.
16/01/21 15:17:54 WARN MetricsSystem: Using default name DAGScheduler for source because spark.app.id
is not set.
Spark context available as sc.
16/01/21 15:17:56 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependenci
es)
16/01/21 15:17:56 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependenci
es)
16/01/21 15:18:00 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.ve
rification is not enabled so recording the schema version 1.2.0
16/01/21 15:18:00 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException
16/01/21 15:18:01 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... usi
ng builtin-java classes where applicable
16/01/21 15:18:02 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependenci
es)
16/01/21 15:18:02 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependenci
es)
16/01/21 15:18:05 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.ve
rification is not enabled so recording the schema version 1.2.0
16/01/21 15:18:05 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException
SQL context available as sqlContext.

scala>
```

Step 3: Run the interactive Apache Spark shell example

1. Test the Spark read-eval-print loop (REPL).
 - a) Create a collection of 20 entries

```
scala> val data = 1 to 20
```

```
scala> val data = 1 to 20
data: scala.collection.immutable.Range.Inclusive = Range(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20)
```

- b) Create a resilient distributed dataset (RDD)

```
scala> val distdata = sc.parallelize(data)
```

```
scala> val distdata = sc.parallelize(data)
distdata: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:23
```

- c) Select and return the even values

```
scala> distdata.filter(_ % 2 == 0).collect()
```

```
scala> distdata.filter(_ % 2 == 0).collect()
res0: Array[Int] = Array(2, 4, 6, 8, 10, 12, 14, 16, 18, 20)
```

2. Exit from the Spark REPL.

```
scala> :q
```

```
scala> :q
Stopping spark context.
bash-4.2# █
```

Step 4: Run the Apache Spark master and worker nodes example

1. Switch to the Spark installation directory.

```
# cd opt/ibm/spark/
```

```
bash-4.2# cd opt/ibm/spark/
bash-4.2# █
```

2. Start the Spark Master server.

```
# ./sbin/start-master.sh
```


```
bash-4.2# ./sbin/start-master.sh
starting org.apache.spark.deploy.master.Master, logging to /opt/ibm/spark/sbin/../logs/spark--org.apac
he.spark.deploy.master.Master-1-422caf422530.out
bash-4.2# █
```

3. Start a Spark Worker in your Docker container and allocate 2GB of memory. Use your Docker container id.

```
# ./sbin/start-slave.sh -m 2G spark://<Docker container id>:7077
```

```
bash-4.2# ./sbin/start-slave.sh -m 2G spark://422caf422530:7077
starting org.apache.spark.deploy.worker.Worker, logging to /opt/ibm/spark/sbin/../logs/spark--org.apac
he.spark.deploy.worker.Worker-1-422caf422530.out
bash-4.2# █
```

4. Using a browser, go to <your Linux ip address>:8080 to verify the Spark GUI is operational.


Spark Master at spark://422caf422530:7077

URL: spark://422caf422530:7077
 REST URL: spark://422caf422530:6066 (cluster mode)
 Alive Workers: 1
 Cores in use: 2 Total, 0 Used
 Memory in use: 2.0 GB Total, 0.0 B Used
 Applications: 0 Running, 0 Completed
 Drivers: 0 Running, 0 Completed
 Status: ALIVE

Worker Id	Address	State	Cores	Memory
worker-20160129205107-172.17.0.10-42871	172.17.0.10:42871	ALIVE	2 (0 Used)	2.0 GB (0.0 B Used)

Running Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

Completed Applications


Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

5. Run an example scala program that estimates the value of Pi. Use your Docker container id.

```
# ./bin/spark-submit --class org.apache.spark.examples.SparkPi --master
spark://<Docker container id>:6066 --deploy-mode cluster /opt/ibm/spark/lib/spark-
examples-1.5.2-hadoop2.6.0.jar
```

```
bash-4.2# ./bin/spark-submit --class org.apache.spark.examples.SparkPi --master spark://422caf422530:6066 --deploy-mode cluster /opt/ibm/spark/lib/spark-examples-1.5.2-hadoop2.6.0.jar
Running Spark using the REST application submission protocol.
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
16/01/29 20:54:03 INFO RestSubmissionClient: Submitting a request to launch an application in spark://422caf422530:6066.
16/01/29 20:54:04 INFO RestSubmissionClient: Submission successfully created as driver-20160129205404-0000. Polling submission state...
16/01/29 20:54:04 INFO RestSubmissionClient: Submitting a request for the status of submission driver-20160129205404-0000 in spark://422caf422530:6066.
16/01/29 20:54:04 INFO RestSubmissionClient: State of driver driver-20160129205404-0000 is now RUNNING.
16/01/29 20:54:04 INFO RestSubmissionClient: Driver is running on worker worker-20160129205107-172.17.0.10-42871 at 172.17.0.10:42871.
16/01/29 20:54:04 INFO RestSubmissionClient: Server responded with CreateSubmissionResponse:
{
  "action": "CreateSubmissionResponse",
  "message": "Driver successfully submitted as driver-20160129205404-0000",
  "serverSparkVersion": "1.5.2",
  "submissionId": "driver-20160129205404-0000",
  "success": true
}
bash-4.2#
```

6. Using a browser, go to <your Linux ip address>:8081 to see the output. Click on the stdout link for your DriverID.


Spark Worker at 172.17.0.10:42871

ID: worker-20160129205107-172.17.0.10-42871
 Master URL: spark://422caf422530:7077
 Cores: 2 (0 Used)
 Memory: 2.0 GB (0.0 B Used)
[Back to Master](#)

Running Executors (0)

ExecutorID	Cores	State	Memory	Job Details	Logs
------------	-------	-------	--------	-------------	------

Finished Executors (1)

ExecutorID	Cores	State	Memory	Job Details	Logs
0	1	KILLED	1024.0 MB	ID: app-20160129205409-0000 Name: Spark Pi User: root	stdout stderr

Finished Drivers (1)

DriverID	Main Class	State	Cores	Memory	Logs	Notes
driver-20160129205404-0000	org.apache.spark.examples.SparkPi	FINISHED	1	1024.0 MB	stdout stderr	

7. Verify the result.



8. Exit the bash shell.

exit

```
bash-4.2# exit
exit
dockerspark2:/data/docker/containers #
```

9. The docker container is now stopped.

docker ps

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
dockerspark2:/data/docker/containers #						

Reference

Links

<https://docs.docker.com/engine/userguide/intro/>

<https://hub.docker.com/r/brunswickheads/spark-1.5.2-s390x/>

<https://spark.apache.org/docs/latest/quick-start.html>